

## Lab 10

```
import numpy as np
import matplotlib.pyplot as plt
import copy
import math

# Function to load data (You need to define your own load_data function or manually provide data)
def load_data():
    # Example dataset
    x_train = np.array([6.1101, 5.5277, 8.5186, 7.0032, 5.8598])
    y_train = np.array([17.592, 9.1302, 13.662, 11.854, 6.8233])
    return x_train, y_train

# Compute cost function
def compute_cost(x, y, w, b):
    m = x.shape[0]
    yp = np.dot(x, w) + b
    total_cost = (1 / (2 * m)) * np.sum((yp - y) ** 2)
    return total_cost

# Compute gradients
def compute_gradient(x, y, w, b):
    m = x.shape[0]
    yp = np.dot(x, w) + b
    dj_dw = (1 / m) * np.sum((yp - y) * x)
    dj_db = (1 / m) * np.sum(yp - y)
    return dj_dw, dj_db

# Gradient descent algorithm
def gradient_descent(x, y, w_in, b_in, cost_function, gradient_function, alpha, num_iters):
    m = len(x)
    J_history = []
    w = copy.deepcopy(w_in)
    b = b_in
    for i in range(num_iters):
        dj_dw, dj_db = gradient_function(x, y, w, b)
        w = w - alpha * dj_dw
        b = b - alpha * dj_db
        if i < 100000:
            cost = cost_function(x, y, w, b)
            J_history.append(cost)
            if i % math.ceil(num_iters / 10) == 0:
                print(f"Iteration {i}: Cost {cost:.2f}")
    return w, b, J_history

# Load data
x_train, y_train = load_data()

# Gradient descent settings
initial_w = 0.0
```

```

initial_b = 0.0
iterations = 1500
alpha = 0.01

# Run gradient descent
w, b, _ = gradient_descent(x_train, y_train, initial_w, initial_b, compute_cost, compute_gradient,
alpha, iterations)
print(f'Optimal parameters: w = {w}, b = {b}')

# Plot the linear fit
m = x_train.shape[0]
predicted = np.zeros(m)
for i in range(m):
    predicted[i] = w * x_train[i] + b
plt.plot(x_train, predicted, c='b')
plt.scatter(x_train, y_train, marker='x', c='r')
plt.title("Profits vs. Population per city")
plt.ylabel('Profit in $10,000')
plt.xlabel('Population of City in 10,000s')
plt.show\(\)

# Predict profit for cities with population of 35,000 and 70,000
predict1 = 3.5 * w + b
predict2 = 7.0 * w + b
print(f'For population = 35,000, we predict a profit of ${predict1 * 10000:.2f}')
print(f'For population = 70,000, we predict a profit of ${predict2 * 10000:.2f}')
[9/9, 4:04 PM] Iman Rafiq PU: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import copy
import math

# Function to load data with error handling
def load_data(file_path):
    try:
        # Load the data without headers
        data = pd.read_csv(file_path, header=None)
        # Assign column names manually
        data.columns = ['Population', 'Profit']
        # Print column names for debugging purposes
        print("Columns in the dataset:", data.columns)
        x_train = np.array(data['Population']) # Population as x
        y_train = np.array(data['Profit'])     # Profit as y
        return x_train, y_train
    except FileNotFoundError:
        print(f'Error: The file {file_path} was not found.')
        raise
    except KeyError as e:

```

```

        print(f"Error: {e}")
        raise
    except pd.errors.EmptyDataError:
        print("Error: The file is empty.")
        raise
    except Exception as e:
        print(f"An unexpected error occurred: {e}")
        raise

# Example usage
file_path = 'ex1data1.txt' # Replace with the path to your actual data file
x_train, y_train = load_data(file_path)

# Compute cost function
def compute_cost(x, y, w, b):
    m = x.shape[0]
    yp = np.dot(x, w) + b
    total_cost = (1 / (2 * m)) * np.sum((yp - y) ** 2)
    return total_cost

# Compute gradients
def compute_gradient(x, y, w, b):
    m = x.shape[0]
    yp = np.dot(x, w) + b
    dj_dw = (1 / m) * np.sum((yp - y) * x)
    dj_db = (1 / m) * np.sum(yp - y)
    return dj_dw, dj_db

# Gradient descent algorithm
def gradient_descent(x, y, w_in, b_in, cost_function, gradient_function, alpha, num_iters):
    m = len(x)
    J_history = []
    w = copy.deepcopy(w_in)
    b = b_in
    for i in range(num_iters):
        dj_dw, dj_db = gradient_function(x, y, w, b)
        w = w - alpha * dj_dw
        b = b - alpha * dj_db
        if i < 100000:
            cost = cost_function(x, y, w, b)
            J_history.append(cost)
            if i % math.ceil(num_iters / 10) == 0:
                print(f"Iteration {i}: Cost {cost:.2f}")
    return w, b, J_history

# Gradient descent settings
initial_w = 0.0
initial_b = 0.0
iterations = 1500
alpha = 0.01

# Run gradient descent

```

```
w, b, _ = gradient_descent(x_train, y_train, initial_w, initial_b, compute_cost, compute_gradient,
alpha, iterations)
print(f"Optimal parameters: w = {w:.2f}, b = {b:.2f}")
```

```
# Plot the linear fit
m = x_train.shape[0]
predicted = np.zeros(m)
for i in range(m):
    predicted[i] = w * x_train[i] + b
plt.plot(x_train, predicted, c="b")
plt.scatter(x_train, y_train, marker='x', c='r')
plt.title("Profits vs. Population per city")
plt.ylabel('Profit in $10,000')
plt.xlabel('Population of City in 10,000s')
plt.show\(\)
```

```
# Predict profit for cities with population of 35,000 and 70,000
predict1 = 3.5 * w + b
predict2 = 7.0 * w + b
print(f'For population = 35,000, we predict a profit of ${predict1 * 10000:.2f}')
print(f'For population = 70,000, we predict a profit of ${predict2 * 10000:.2f}')
```